


Accuracy vs. training compute: dense and sparse cost the same FLOPs at a fixed epoch count (unstructured pruning skips no compute) — SAM's 2x-per-step cost is the only real compute asymmetry in the whole comparison

ResNet18 (sparse: $s = 0.999$)



VGG16 (sparse: $s = 0.9995$)

